CSE141L Lab 3 Caching Optimizations Worksheet1

Name: _____ Student ID: _____

Instructions

- Complete this worksheet while reading/working through the lab write up. The worksheet doesn't make sense without the lab.

- The point values are listed for each question. Altering the size of the cells will cost you 1 point. The write up portion of the lab is 30% of your total point for the lab as shown in the lab's README.md

Cache and dataset characteristics

P1 (4pt) Find out the dimensions (number of data elements) of the following tensors/vectors used in `fc_layer_t::activate` for the cifar100 dataset and fill the following table

Tensor/Vector	Number of Data Elements
`in`	_____
`out`	_____
`weights`	_____
`activation_input`	_____

P2 (4pt) Calculate the size (in Bytes) of the following tensors/vectors used in `fc_layer_t::activate` for the cifar100 dataset and fill the following table

Tensor/Vector	Size in Bytes
`in`	_____
`out`	_____
`weights`	_____
`activation_input`	_____

P3 (4pt) How much of each of these data structures used in `fc_layer_t::activate()` will fit in the L1 and L2 cache (Note The cache size for our machine is L1-dCache: 32kb, L1-Icache: 32kb, L2: 256kb, L3: 8Mb)?

tensor	% that'll fit in L1	% that'll fit in L2
`in`	_____	_____
`out`	_____	_____
`weights`	_____	_____
`activation_input`	_____	_____

Understanding Tensor_t

Given `tensor_t<double> foo(tdsize(4,3,5,7))`, answer the following (double are 8 bytes)(Hint: Look at lecture slides) :

P1 (1pt) How many elements are there in `foo`?

P2 (1pt) What's the linear index of element (1,1,1,1)?

P3 (1pt) How far apart are elements that differ by 1 in each dimension?

dim.	distance in bytes	distance in linear index
x		
y		
z		
b		

Tier 1: Reordering and Tiling loops in `fc_layer_t::activate`

P1(a) (3pt) Fill out the following table. Report the Misses per Instruction by using the performance counters (there should be a column for "MPI" in the reported data when running with L1/2/3.cfg)

Cache-Level	Miss rate - Base	Miss rate - loop reordering	Miss rate - Tiling
L1			
L2			
L3			

P1(b) (1pt) Were there any differences in the miss rate observed using the performance counters and moneta? What could contribute to the differences? (A brief answer is fine)

```
   Your answer here
```

P2 (4pt) Change the order of loops from `b i n` to `b n i` in `fc_layer_t::activate` and report the speedup.

Speedup after loop reordering : _____

P3 (4pt) Block the loop `n` in `fc_layer_t::activate` with the tile sizes 1, 2, 4, 8, 16 and fill out the table below.

Dataset	Step size	Blocked implementation time	Speedup vs step size == 1
cifar100	1	_____	_____
cifar100	2	_____	_____
cifar100	4	_____	_____
cifar100	8	_____	_____
cifar100	16	_____	_____

P4 (4pt) In a single line graph, plot the speed up against the different block sizes for blocking the loop `n` in `fc_layer_t::activate` . Block size is the independent vairable.

```
    Your graph here
```

P5 (4pt) Consider the blocksize which gave maximum speedup in the previous question P4 and fill out the following table

1. Base implementation time : _____
2. Implementation time of your optimized solution : _____
3. Base implementation L1 misses : _____
4. Your fastest solution L1 misses : _____

P6 (3pt) Insert the memory access patterns (take screenshots from moneta) for loop orders b-i-n, b-n-i and nn-b-n-i. Do this for weights and pass the runtime options that set scale to 4 and reps to 1. The dataset should be cifar100. Leave the cache lines and block size fields as they are but set the max accesses to around 2 million. In the starter code, there is an example of where and what to put to run the moneta please take a look at those).

```
    memory access pattern with loop order b-i-n
```

```
    memory access pattern with loop order b-n-i
```

memory access pattern with loop order nn-b-n-i